Beyond Hate: Differentiating Incivil and Intolerant Speech in Multimodal Content Moderation

Nils A. Herrmann[1], **Jingyi He**[1],

[1]Technical University of Munich
nils_hermann@outlook.de, holly.he@tum.de

Abstract

Current multimodal toxic speech detection models usually classify content as either hateful or not using a single binary label. However, this coarse approach confuses two fundamentally different characteristics of expression: incivility and intolerance. Building on communication science research, we propose a fine-grained annotation scheme that distinguishes between *tone* (incivility) and *content* (intolerance). We empirically test the advantages of this distinction by re-annotating a subset of the Hateful Memes dataset and fine-tuning a vision-language model (LLaVA-1.6-7B) under different training configurations. Our results show that using our fine-grained annotation scheme to label a small subset of the dataset improves model performance. Furthermore, we uncover systematic misclassification biases, such as under-detection of civil but intolerant content, that highlight the importance of moving beyond single hatefulness labels. This work contributes to data-centric AI, which improves the reliability and accuracy of moderation systems by enhancing data quality.

Introduction

Informal conversations shape people's understanding of the world and public issues in democratic societies (Conover and Searing 2005). In our digital world, these conversations take various forms including text-image combinations or "memes" (Milner 2018). Although memes are usually harmless and humorous, they can be dangerous if they promote extremist views (Askanius 2021). In the fast-paced environment of the internet, quickly identifying toxic memes is crucial to preventing the spread of harmful content.

Since the Hateful Meme Detection Challenge in 2020 (Kiela et al. 2020), the detection of memes has received particular attention (Pandiani, Sang, and Ceolin 2024). Various computational approaches have been suggested for identifying toxic memes. These include *unimodal* (Waseem et al. 2017), *multimodal* (Hossain, Hoque, and Hossain 2023; Liang et al. 2022; Kumar and Nandakumar 2022; Qu et al. 2023; Koutlis, Schinas, and Papadopoulos 2023) and *generative models* (Lin et al. 2023, 2024; Van and Wu 2023; Cao et al. 2022; Jain et al. 2023).

Toxic speech is often operationalized as either hateful or not, which oversimplifies the complexity of online discourse. This coarse *single binary label* approach can mix up *incivil* but potentially constructive speech with genuinely *intolerant* and harmful content. From a normative perspective, this undermines democratic dialogue by suppressing incivil but legitimate expressions of disagreement, such as emotionally charged criticism. Incivil expressions may still contribute to public debate despite their tone (Rossini 2022). In contrast, intolerant speech explicitly seeks to dehumanize or exclude individuals based on identity and should be treated differently. From a practical perspective, using a fine-grained *multi-label approach* which distinguishes between tone (incivility) and content (intolerance) enables the development of nuanced annotation schemes, which in turn can lead to more reliable training data. According to research in data-centric AI (Zha et al. 2023), improvements in model performance often come from better data rather than from model tuning.

This paper contributes to data-centric AI by introducing a fine-grained multi-label operationalization of hate speech that separates incivility from intolerance. Based on this definition, we develop a new annotation scheme that captures this distinction and apply it to the widely used Hateful Memes dataset. Finally, we empirically test how the fine-grained annotation impacts the performance of a multimodal hate speech detection model. Our approach not only supports democratic values by protecting legitimate disagreement but also helps build more accurate and interpretable detection systems.

Related Work

Current Multimodal Computational Approaches

Effectively identifying hateful memes requires multimodal understanding, as the meaning often emerges from the interaction between visual and textual components. Unimodal models which process only text or image, struggle to capture nuanced semantics, particularly in the presence of confounders where hatefulness arises from the interplay of modalities (Kiela et al. 2020).

Two major computational paradigms have emerged for this task. The first is the *specialist approach*, which fine-tunes pretrained vision-language encoders using task-specific fusion mechanisms. These models are tailored exclusively for hateful meme detection. Notable examples include HateCLIPper, which aligns CLIP embeddings through cross-modal interaction modules (Kumar and Nandakumar

2022), and MemeCLIP, which introduces feature adapters and a cosine classifier to address overfitting and class imbalance (Shah et al. 2024). The second is the *generalist approach*, which employs large multimodal models with broad vision-language capabilities. These include off-the-shelf models such as LLaVA (Liu et al. 2023) and Qwen2VL (Wang et al. 2024), as well as model adaptation strategies like RA-HMD (Mei et al. 2025), which introduces a two-stage fine-tuning framework and retrieval-augmented classification. While generalist models like LLaVA were not designed specifically for hate detection, their capacity to generate textual rationales and generalize across tasks makes them especially promising for evaluating nuanced definitions of hatefulnes. Recent benchmark results reported in Mei et al. (2025) support this potential. On the HatefulMemes dataset, supervised fine tuning (SFT) yields an accuracy of approximately 78% for both LLaVA and Qwen2VL. With architectural enhancements and retrieval integration, RA-HMD fine tuned Qwen2VL achieves a new state of the art accuracy of 82.1%. In comparison, specialist models based on CLIP achieve lower performance: the original CLIP model reaches 72% accuracy (Radford et al. 2021), and HateCLIPper improves this to 76% (Kumar and Nandakumar 2022). These findings highlight the strength of generalist models when appropriately adapted for the task. Among widely adopted off the shelf models, LLaVA and Qwen2VL currently offer the best balance of accessibility and performance.

Multi-label Hatefulness

There are recent efforts to disentangle different forms of harmful online expression by distinguishing between incivility and intolerance. Rossini (2022) argues that these two forms of speech are not only conceptually distinct but also empirically separable in political discourse. Incivility, marked by rudeness, often emerges in discussions that reflect the vibrant, even confrontational, nature of democratic deliberation. By contrast, intolerance refers to speech that undermines democratic values by attacking rights, identities, or pluralism. Rossini (2022) provides evidence that incivility can occur even in productive political engagement, while intolerance appears in contexts that exacerbate harm. It is crucial to recognise this distinction in order to reframe debates about moderation, which differentiate between an uncivil tone and the more substantial threats posed by intolerant content.

Our approach builds on Bianchi et al. (2022), who apply a multi-label framework to a large corpus of immigration-related tweets from the US and UK. They use text-based models to detect both incivility and intolerance. This multi-label approach improves understanding of harmful discourse, producing models that outperform those trained on coarse-grained datasets. These results reinforce the idea that multi-label schemes can improve both the conceptual clarity and the empirical performance of content moderation.

Data

Hateful Meme Dataset

The Hateful Memes dataset (Kiela et al. 2020) was introduced as a technical benchmark for multimodal hate speech detection. Its goal is to test whether models can effectively combine image and text information to detect hateful content. To make the task more challenging and avoid shortcuts based on unimodal cues, the dataset includes "benign confounders", examples where neither the image nor the text alone is sufficient to determine hatefulness.

The dataset contains approximately 10,000 memes, each consisting of an image paired with embedded text. Labels are binary: memes are either marked as hateful or not. The challenge defines hateful content as follows:

> A direct or indirect attack on people based on characteristics, including ethnicity, race, nationality, immigration status, religion, caste, sex, gender identity, sexual orientation, and disability or disease. We define attack as violent or dehumanizing (comparing people to non-human things, e.g. animals) speech, statements of inferiority, and calls for exclusion or segregation. Mocking hate crime is also considered hate speech.

Although previous research (Davidson et al. 2017) proposed a three-way distinction separating hateful, offensive, and normal content, the Hateful Memes challenge opted for a single binary label. They justify this choice by arguing that the distinction between hatefulness and offensiveness is "murky, and the classification decision less actionable" (Kiela et al. 2020, p. 3).

A single binary label is appropriate in the context of a technical benchmark, since the goal in such settings is to simplify the task, standardize evaluation, and focus on model performance under controlled conditions. Challenge settings often prioritize clean comparisons over realism and tend to abstract away many of the complexities found in actual user-generated content. However, real-world scenarios demand a more fine-grained approach that can capture the diverse and ambiguous nature of online speech. In such contexts, we argue that distinguishing between incivility and intolerance is essential.

From a normative perspective, we follow communication science research that emphasizes the conceptual and empirical separability of these two forms of speech in political discourse (Coe, Kenski, and Rains 2014; Kosmidis and Theocharis 2020; Rossini 2022; Kümpel and Unkel 2023). From a practical point of view, failing to model incivility can cause problems. It creates conceptual ambiguity, making it harder to determine whether speech is genuinely intolerant or simply confrontational in tone. This blurring of boundaries can lead to misclassifications, such as flagging civil disagreement as hate or overlooking polite but harmful statements. The study by Kern et al. (2023) shows that annotation clarity and agreement are highest when annotators classify tone (offensive language) and content (hate speech) simultaneously, underscoring the value of distinguishing these dimensions. Finally, omitting incivility deprives models of

rich linguistic cues that can reveal cases where tone amplifies or conceals harmful content.

Annotation of the Hateful Memes Dataset

To introduce a more nuanced understanding of hatefulness, we annotate a subset of the Hateful Memes dataset with two variables that distinguish between incivility and intolerance. Specifically, we annotate 6.75% of the total dataset, corresponding to approximately 500 memes. Our annotation scheme is grounded in definitions from communication science, drawing particularly on the work of Rossini (2022) and Coe, Kenski, and Rains (2014). Incivility and intolerance are treated as two separate variables, each capturing a different dimension of discourse.

Incivility is defined as

> Rude, disrespectful or dismissive *tone* towards others as well as opinions expressed with antinormative intensity.

Intolerance is defined as

> *Behaviors* that are threatening to democracy and pluralism—such as prejudice, segregation, hateful or violent speech, and the use of stereotyping in order to disqualify others and groups.

Annotators are provided with definitions of incivility and intolerance, along with lists of their respective sub-types. The scheme specifies three types of incivility and nine types of intolerance (see Appendix). Annotators must indicate the specific sub-type whenever either is present. For fine-tuning and evaluation, the annotations are consolidated into two separate binary labels.

Table 1 contrasts the two annotation strategies used in this study. The original approach applies a coarse single-label scheme, classifying content as either *Hateful* or *Neutral*. Our approach adopts a fine-grained multi-label scheme with two binary labels, distinguishing *Intolerant* from *Neutral* and *Incivil* from *Neutral*.

Annotation type	Labels	% of dataset annotated
Coarse single-label	Hateful vs. Neutral	100%
Fine-grained multi-label	Intolerant vs. Neutral Incivil vs. Neutral	6.75%

Table 1: Two annotation schemes used in this study, with their respective coverage in the dataset.

Empirical Methodology

Toxicity Identification

In hateful meme detection, harmfulness often emerges not from the text or image alone but from their interaction, requiring models to perform joint multimodal reasoning rather than isolated analysis.

To effectively leverage from both text and image information, we adopt Vision-Language Model (VLM) which is capable of learning from both visual and textual inputs in a unified manner. Specifically, we utilize **LLaVA-1.6-Mistral-7B**, a recent model that integrates a pretrained image encoder (e.g., CLIP ViT) with a large language model (e.g., Mistral) via a multimodal projection mechanism. The visual input is processed into a sequence of patch embeddings, which are projected into the same embedding space as the language model tokens. These visual tokens are then inserted into the language model's input stream, allowing the model to attend jointly over both modalities from the earliest layers. Figure 1 shows our whole pipeline.

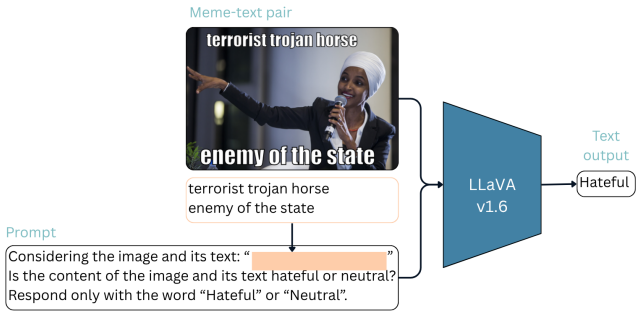

Figure 1: Evaluation pipeline. A meme–text pair is passed to the LLaVA v1.6 model along with a prompt asking whether the content is "Hateful" or "Neutral". The model generates a binary classification based on both visual and textual input.

This design enables **cross-modal information flow** throughout the transformer layers and supports fine-grained alignment between image regions and linguistic context. LLaVA's unified token sequence allows it to reason about visual-textual relationships. Such architecture is particularly well-suited to the toxic content detection task, where subtle inferences must be drawn from the combination of visual cues and textual phrasing.

Summarized, we fine-tune LLaVA-1.6 on the hateful meme dataset to evaluate its ability to capture these multimodal signals and correctly classify instances where toxicity emerges from the image-text pairing, rather than either modality alone.

Cross-Granularity Transfer Learning of Hatefulness Understanding

We adopt a dual-phase experimental setup motivated by transfer learning and representation generalization paradigms. In addition to standard training, where models are trained and evaluated on the same label space, we explore how varying levels of annotation granularity impact the model's internal representations and generalization behavior.

To examine the transferability and robustness of learned representations, we define two cross-granularity learning configurations:

- **Coarse-to-Fine Evaluation:** We train a model on the full single-label hatefulness dataset and evaluate it on the multi-label subset to test whether coarse-grained super-

vision can implicitly capture and generalize to the finer-grained categories of incivility and intolerance.

- **Fine-to-Coarse Evaluation:** We train a model solely on the multi-label subset and evaluate it on the full single-label dataset. This setup mirrors a few-shot domain-adaptive fine-tuning paradigm, where a small amount of semantically rich supervision is used to guide the learning of broader hatefulness concepts. The goal is to examine whether training on diverse expressions of toxicity equips the model with generalized representations that can effectively perform binary classification on a wider set of unlabeled memes.

Joint Learning with Heterogeneous Label Granularity

Besides the cross-granularity evaluation, we propose a joint-learning setup that the model leverages both coarse and fine-grained annotations simultaneously. This reflects the real-life scenarios in which rich semantic supervision is only available for a small subset of samples, while the majority of the dataset receive only coarse annotations.

More concretely, samples from the full dataset contribute to a binary hatefulness classification loss, while samples from the manually labeled subset additionally supervise two auxiliary prediction heads corresponding to uncivil and intolerant subcategories. The shared vision-language encoder is trained jointly across all available signals, encouraging the development of representations that are both broadly discriminative and semantically aware.

This partially supervised multi-task learning strategy allows the model to:

- Utilize the entire dataset during training without discarding coarsely labeled samples.
- Enhance interpretability by grounding predictions in interpretable subcategories where available.
- Implicitly regularize the encoder by learning from both coarse and fine-grained supervision, potentially improving generalization on downstream tasks.

Performance Biases

To understand performance biases in hatefulness prediction, we conduct a false positive and false negative analysis using contingency tables. Specifically, we examine whether the baseline model's misclassifications are associated with the true labels for incivility and intolerance, thereby revealing potential biases in its predictions. To statistically assess the presence and significance of these biases, we employ a Fisher exact test. Intuitively, this approach evaluates whether certain types of memes, such as incivil but non-hateful or civil but hateful, are disproportionately misclassified relative to their counterparts.

Relationship Incivility/Intolerance and Hatefulness

To examine how the commonly used single-label conceptualization of toxicity aligns with a multi-label approach, we analyze the relationship between the original hatefulness labels in the Hateful Memes dataset and our incivility and intolerance annotations. We examine two competing hypotheses:

- **H1:** A meme is labeled as hateful if it is *incivil* OR *intolerant*.
- **H2:** A meme is labeled as hateful if it is *intolerant*.

To test these hypotheses, we compute the correlation between the original single hatefulness labels and our multi-label annotations for incivility and intolerance. Additionally, we conduct chi-squared tests of independence to assess the statistical association between hatefulness and the presence of incivility and intolerance. This analysis helps clarify whether the binary notion of hatefulness more closely tracks with intolerance alone, or with a broader set of norm violations that includes incivility.

Results

Toxicity Identification

The results of the experiment are presented in Table 2.

As outlined in the Method section, we evaluate three experimental configurations: **baseline, transfer learning,** and **joint learning**. Each of these settings highlights different aspects of how the model processes coarse and fine-grained hatefulness annotations.

In the **transfer learning** setting, we observe that a model trained exclusively on single hatefulness labels performs poorly when evaluated on the fine-grained subset. The accuracies of 66.7% and 65.7% across the two configurations indicate only minor improvements over random guessing. This finding suggests that, despite being trained on a large dataset for general hatefulness detection, the model **lacks the capacity to capture more subtle and nuanced forms of hateful content**, such as incivility or intolerance. This limitation highlights the gap between coarse supervision and the complex nature of real-world hateful expressions.

In contrast, the **joint learning** experiments provide encouraging evidence that introducing even a small amount of semantically rich supervision can lead to notable improvements. By incorporating approximately 600 manually annotated fine-grained samples, which account for only **6.75%** of the entire dataset, the model achieves a **4%** accuracy improvement in single-label hatefulness prediction compared to the baseline. This result indicates that fine-grained annotations not only enhance the interpretability of the model's internal representations but also provide a form of regularization that **strengthens its overall generalization performance**. The improvement achieved with such a small fraction of detailed annotations suggests that combining broad, coarse-labeled data with strategically selected fine-grained labels is a promising strategy for advancing hatefulness detection.

Analysis	Finetune Split	Evaluation Split	Accuracy	F1
Baseline	Hateful	Hateful	0.741	0.736
	Incivil/Intolerant^{600}	Incivil^{600}	0.770	0.769
		Intolerant^{600}	**0.775**	**0.770**
Transfer Learning	Incivil/Intolerant^{600}	Hateful	0.650 (-0.12)	0.635 (-0.14)
	Hateful	Incivil^{600}	0.667 (-0.13)	0.633 (-0.18)
		Intolerant^{600}	0.657 (-0.15)	0.652 (-0.15)
Joint Learning	Incivil/Intolerant^{600}	Hateful	**0.768** (+0.04)	**0.758** (+0.03)
	+	Incivil^{600}	**0.789** (+0.02)	**0.821** (+0.07)
	Hateful	Intolerant^{600}	0.716 (-0.08)	0.733 (-0.05)

Table 2: Performance across different training and evaluation setups. The best results per label are shown in bold. Deltas are shown in parentheses relative to the corresponding baseline. Green indicates improvement and red indicates decline.

Performance Biases

Table 3 summarizes the biases of the baseline model when predicting hatefulness.

Subset (Misclassified)	Error Type	Error Likelihood
Incivil non-hateful memes	False positive	1.2x
Civil hateful memes	False negative	2.0x**
Intolerant non-hateful memes	False positive	2.2x**
Tolerant hateful memes	False negative	1.1x

Table 3: Hatefulness misclassification likelihoods for different subsets. Asterisks ** indicate statistical significance based on Fisher's exact test ($\alpha = 0.05$).

The results show that misclassifications are influenced by the lack of distinction between incivility and intolerance. This introduces bias in how the model predicts hatefulness. Two statistically significant associations are particularly important. Civil but hateful memes are twice as likely to be missed, suggesting the model often overlooks harmful content when it appears polite. Intolerant but non-hateful memes are 2.2 times more likely to be wrongly flagged as hateful, indicating that signs of intolerance increase the chance of over-flagging.

These findings highlight the importance of separating tone from content. Without this distinction, moderation models risk missing actual harm and falsely censoring content that does not meet the threshold for hatefulness. Recognizing these nuances is essential for reducing bias and improving prediction accuracy.

Relationship Incivility/Intolerance and Hatefulness

The results in Table 4 suggest that hatefulness in the Hateful Memes dataset aligns more closely with intolerance than with the broader condition of incivility or intolerance. The match rate[1] is higher for the intolerance-only hypothesis (87.9% vs. 81.1%). Likewise, the correlation between hatefulness and intolerance ($r = 0.757$) is stronger than with the combined incivility-or-intolerance condition ($r = 0.655$). In both cases, chi-squared tests confirm statistically significant associations at $\alpha = 0.001$.

	Hateful = Intolerant	Hateful = Incivil OR Intolerant
Perfect Matches	531/604 (87.9%)	490/604 (81.1%)
Correlation	0.757	0.655
χ^2 (α=0.001)	Sig.	Sig.

Table 4: Comparison of hatefulness labels with two conceptualizations. Correlation is the phi coefficient (equivalent to Pearson's r for binary variables). χ^2 refers to the chi-squared test of independence; "Sig." indicates statistical significance at $\alpha = 0.001$.

Conclusion and Future Research

This study shows that coarse binary labels for hatefulness constrain model understanding and obscure critical differences between forms of toxic speech. By introducing fine-grained annotations for incivility and intolerance, grounded in social science research, we enable models to disentangle the tonal and substantive dimensions of toxicity. Even with only 6.75% of the dataset manually annotated at this level, joint learning with both coarse and fine-grained labels improves LLaVA-1.6-7B's performance and interpretability. This demonstrates the value of strategically integrating rich supervision into large-scale datasets.

Our results further reveal that the Hateful Memes dataset's operationalization of toxicity aligns more closely with intolerance than with a broader conception that includes incivility. This bias manifests in systematic errors, such as overlooking civil but hateful content and over-flagging intolerant yet non-hateful memes. These findings highlight the importance of explicitly modeling tone to avoid both under-detection of toxic speech and over-moderation of permissible discourse.

From a methodological standpoint, our work underscores the promise of hybrid annotation strategies that combine

[1] Proportion of cases where the label predicted under a given hypothesis matches the original hatefulness label in the dataset.

large-scale coarse labels with targeted fine-grained annotations to achieve gains in both accuracy and transparency. From a normative perspective, separating incivility from intolerance supports more democratic moderation practices by protecting legitimate disagreement while prioritizing toxic speech detection that threatens rights, dignity, or inclusion.

Future research should test the multi-label approach on real-world datasets, with a particular focus on validating bias analysis for marginalised sub-groups. Results should also be validated using OCR-derived text instead of embedded meme text to simulate deployment conditions. Another promising direction is to test the multi-label approach with specialised CLIP-based models. In the future, state-of-the-art performance in toxicity detection could come from a more data-centric approach. This would involve expanding the label set and using joint learning to combine coarse and fine-grained supervision.

Ethical Considerations

In line with the European Union's guidelines on Trustworthy AI (Commission 2019), this project aims to mitigate harm by detecting hateful content. We follow a data-centric approach that enhances model detection capabilities and increases transparency through the introduction of a detailed annotation codebook. This codebook provides nuanced definitions of incivility and intolerance, supporting more precise and interpretable moderation decisions. Several ethical considerations and limitations remain. First, striking a balance between preventing harm and preserving freedom of speech is a central concern. Overly restrictive moderation risks suppressing legitimate discourse, while insufficient intervention can allow harmful content to persist. Second, assessments of content are not entirely objective; they are often shaped by the annotator's individual perspective and sociocultural context (Hettiachchi et al. 2023), which can introduce variability and bias into annotations. Third, adaptability is a major challenge. Internet memes are fast-evolving communicative artifacts that reflect rapidly shifting cultural symbols, formats, and narratives. To remain effective, moderation systems must be resilient while also capable of adapting to emerging patterns in both form and meaning.

References

Askanius, T. 2021. On Frogs, Monkeys, and Execution Memes: Exploring the Humor-Hate Nexus at the Intersection of Neo-Nazi and Alt-Right Movements in Sweden. 22(2): 147–165.

Bianchi, F.; Hills, S. A.; Rossini, P.; Hovy, D.; Tromble, R.; and Tintarev, N. 2022. "It's Not Just Hate": A Multi-Dimensional Perspective on Detecting Harmful Speech Online. .

Cao, R.; Lee, R. K.-W.; Chong, W.-H.; and Jiang, J. 2022. Prompting for Multimodal Hateful Meme Classification. In Goldberg, Y.; Kozareva, Z.; and Zhang, Y., eds., *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, 321–332. Association for Computational Linguistics.

Coe, K.; Kenski, K.; and Rains, S. A. 2014. Online and Uncivil? Patterns and Determinants of Incivility in Newspaper Website Comments. 64(4): 658–679.

Commission, E. 2019. *Ethics Guidelines for Trustworthy AI*. Publications Office of the European Union. ISBN 978-92-76-11998-2.

Conover, P. J.; and Searing, D. D. 2005. Studying 'Everyday Political Talk' in the Deliberative System. 40(3): 269–283.

Davidson, T.; Warmsley, D.; Macy, M.; and Weber, I. 2017. Automated Hate Speech Detection and the Problem of Offensive Language. .

Hettiachchi, D.; Holcombe-James, I.; Livingstone, S.; family=Silva, p.-u., given=Anjalee; Lease, M.; Salim, F. D.; and Sanderson, M. 2023. How Crowd Worker Factors Influence Subjective Annotations: A Study of Tagging Misogynistic Hate Speech in Tweets. 11: 38–50.

Hossain, E.; Hoque, M. M.; and Hossain, M. A. 2023. An Inter-modal Attention Framework for Multimodal Offense Detection. In Vasant, P.; Weber, G.-W.; Marmolejo-Saucedo, J. A.; Munapo, E.; and Thomas, J. J., eds., *Intelligent Computing & Optimization*, 853–862. Springer International Publishing. ISBN 978-3-031-19958-5.

Jain, R.; Maity, K.; Jha, P.; and Saha, S. 2023. Generative Models vs Discriminative Models: Which Performs Better in Detecting Cyberbullying in Memes? In *2023 International Joint Conference on Neural Networks (IJCNN)*, 1–8.

Kern, C.; Eckman, S.; Beck, J.; Chew, R.; Ma, B.; and Kreuter, F. 2023. Annotation Sensitivity: Training Data Collection Methods Affect Model Performance. In Bouamor, H.; Pino, J.; and Bali, K., eds., *Findings of the Association for Computational Linguistics: EMNLP 2023*, 14874–14886. Association for Computational Linguistics.

Kiela, D.; Firooz, H.; Mohan, A.; Goswami, V.; Singh, A.; Ringshia, P.; and Testuggine, D. 2020. The Hateful Memes Challenge: Detecting Hate Speech in Multimodal Memes. In *Advances in Neural Information Processing Systems*, volume 33, 2611–2624. Curran Associates, Inc.

Kosmidis, S.; and Theocharis, Y. 2020. Can Social Media Incivility Induce Enthusiasm?: Evidence from Survey Experiments. 84(S1): 284–308.

Koutlis, C.; Schinas, M.; and Papadopoulos, S. 2023. MemeFier: Dual-stage Modality Fusion for Image Meme Classification. .

Kumar, G. K.; and Nandakumar, K. 2022. Hate-CLIPper: Multimodal Hateful Meme Classification Based on Cross-modal Interaction of CLIP Features. .

Kümpel, A. S.; and Unkel, J. 2023. Differential Perceptions of and Reactions to Incivil and Intolerant User Comments. 28(4): zmad018.

Liang, X.; Huang, Y.; Liu, W.; Zhu, H.; Liang, Z.; and Chen, L. 2022. TRICAN: Multi-Modal Hateful Memes Detection with Triplet-Relation Information Cross-Attention Network. In *2022 International Joint Conference on Neural Networks (IJCNN)*, 1–8.

Lin, H.; Luo, Z.; Gao, W.; Ma, J.; Wang, B.; and Yang, R. 2024. Towards Explainable Harmful Meme Detection

through Multimodal Debate between Large Language Models. .

Lin, H.; Luo, Z.; Ma, J.; and Chen, L. 2023. Beneath the Surface: Unveiling Harmful Memes with Multimodal Reasoning Distilled from Large Language Models. In Bouamor, H.; Pino, J.; and Bali, K., eds., *Findings of the Association for Computational Linguistics: EMNLP 2023*, 9114–9128. Association for Computational Linguistics.

Liu, H.; Li, C.; Wu, Q.; and Lee, Y. J. 2023. Visual Instruction Tuning. 36: 34892–34916.

Mei, J.; Chen, J.; Yang, G.; Lin, W.; and Byrne, B. 2025. Robust Adaptation of Large Multimodal Models for Retrieval Augmented Hateful Meme Detection. .

Milner, R. M. 2018. *The World Made Meme: Public Conversations and Participatory Media*. MIT Press. ISBN 978-0-262-53522-9.

Pandiani, D. S. M.; Sang, E. T. K.; and Ceolin, D. 2024. Toxic Memes: A Survey of Computational Perspectives on the Detection and Explanation of Meme Toxicities. .

Qu, Y.; He, X.; Pierson, S.; Backes, M.; Zhang, Y.; and Zannettou, S. 2023. On the Evolution of (Hateful) Memes by Means of Multimodal Contrastive Learning. 293–310. IEEE Computer Society. ISBN 978-1-66549-336-9.

Radford, A.; Kim, J. W.; Hallacy, C.; Ramesh, A.; Goh, G.; Agarwal, S.; Sastry, G.; Askell, A.; Mishkin, P.; Clark, J.; Krueger, G.; and Sutskever, I. 2021. Learning Transferable Visual Models From Natural Language Supervision. .

Rossini, P. 2022. Beyond Incivility: Understanding Patterns of Uncivil and Intolerant Discourse in Online Political Talk. 49(3): 399–425.

Shah, S. B.; Shiwakoti, S.; Chaudhary, M.; and Wang, H. 2024. MemeCLIP: Leveraging CLIP Representations for Multimodal Meme Classification. In Al-Onaizan, Y.; Bansal, M.; and Chen, Y.-N., eds., *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 17320–17332. Association for Computational Linguistics.

Van, M.-H.; and Wu, X. 2023. Detecting and Correcting Hate Speech in Multimodal Memes with Large Visual Language Model. .

Wang, P.; Bai, S.; Tan, S.; Wang, S.; Fan, Z.; Bai, J.; Chen, K.; Liu, X.; Wang, J.; Ge, W.; Fan, Y.; Dang, K.; Du, M.; Ren, X.; Men, R.; Liu, D.; Zhou, C.; Zhou, J.; and Lin, J. 2024. Qwen2-VL: Enhancing Vision-Language Model's Perception of the World at Any Resolution. .

Waseem, Z.; Davidson, T.; Warmsley, D.; and Weber, I. 2017. Understanding Abuse: A Typology of Abusive Language Detection Subtasks. In Waseem, Z.; Chung, W. H. K.; Hovy, D.; and Tetreault, J., eds., *Proceedings of the First Workshop on Abusive Language Online*, 78–84. Association for Computational Linguistics.

Zha, D.; Bhat, Z. P.; Lai, K.-H.; Yang, F.; Jiang, Z.; Zhong, S.; and Hu, X. 2023. Data-Centric Artificial Intelligence: A Survey. .

Appendix

Annotation Scheme

This appendix provides an overview of the sub-types used in our annotation scheme for incivility and intolerance. Annotators labeled each meme using the definitons below.

Incivility Sub-types

To guide the annotation of incivility, we use a typology that captures different forms of uncivil expression. Each subtype is described in both text and image, since the idea of "tone" can appear differently in each. Textual incivility may involve word choice or phrasing, while visual incivility can appear through imagery or symbols. Table 5 lists the categories and provides a description for each modality.

Nr.	Form of Incivility	Text	Visual
0	Civil	Absence of incivility.	Absence of offensive, hostile, or degrading imagery.
1	Profane or vulgar	Use of explicit profanities or vulgarities, regardless of whether they are directed at a person or entity.	Images containing profanity, obscene gestures, or crude/vulgar symbolism. Sexual or violent images aimed at intimidating, shocking, or degrading public discourse.
2	Attacks	Derogatory or pejorative language directed at specific individuals or groups. These can focus on personal characteristics (e.g., appearance), traits, character, or choices.	Images targeting identifiable individuals or demographic groups with pejorative, mocking, or dehumanising intent.
3	Aspersions	Attacks targeting groups, organizations, or institutions aiming to undermine credibility, legitimacy, or moral standing.	Images delegitimising institutions, policies, or parties through ridicule, slurs, or hostile visual metaphors.

Table 5: Subtypes of Incivility in Text and Visual Content. Used in the annotation scheme.

Intolerance Sub-types

To guide the annotation of intolerance, we use a typology that captures different forms of harmful content. Unlike incivility, which can manifest in different ways in text and images, the meaning of intolerant messages tends to remain consistent across modalities. Table 6 presents each subtype alongside a brief description.

Nr.	Form of Intolerance	Description
0	Tolerant / Neutral	Absence of expressions of intolerance.
1	Threats to Individual Rights	Claims that certain groups should not have equal civil, political, or human rights.
2	Intolerance Toward Political Positions	Delegitimizing opposing political views or calling for elimination of political actors.
3	Racism	Discriminatory, stereotypical, hateful, or prejudicial speech toward racial minorities.
4	Social/Economic Intolerance	Discriminatory, stereotypical, hateful, or prejudicial speech toward others based on education level, social status, or income.
5	Gender and Sexual Intolerance	Discriminatory, stereotypical, hateful, or prejudicial speech toward women and/or LGBTQ+ individuals based on gender status or sexual orientation.
6	Religious Intolerance	Discriminatory, stereotypical, hateful, or prejudicial speech toward individuals or groups based on religion.
7	Offensive Stereotyping	Cultural, regional, physical, or professional stereotyping with derogatory framing.
8	Violent Threats	Calls to harm individuals or institutions, or support for violence.
9	Ableism	Prejudice, stereotypes, mockery, exclusion, or hostility toward people with physical, intellectual, sensory, or mental disabilities.

Table 6: Subtypes of Intolerance in Text Content. Used in the annotation scheme.